

'CprE 381: Computer Organization and Assembly-Level Programming

Project Part 1 Report

Team Members: Jack Sebahar

Peter Wissman

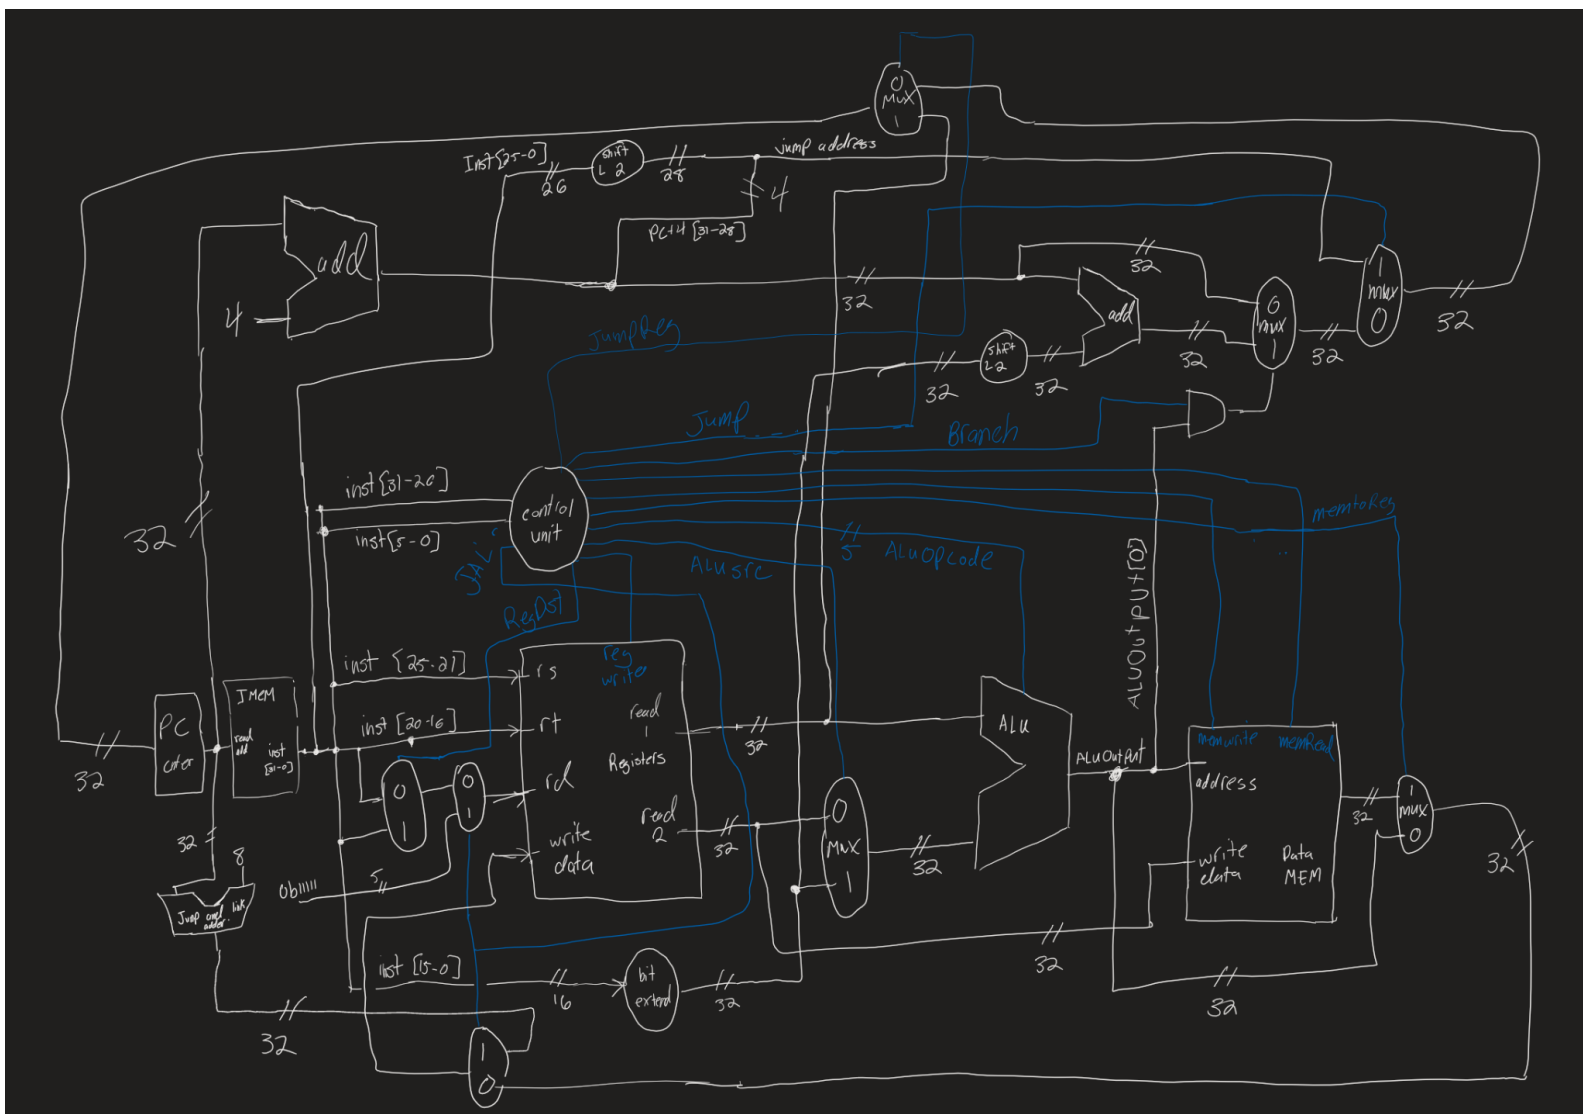
John Lavigne

Project Teams Group #: 1

Refer to the highlighted language in the project 1 instruction for the context of the following questions.

[Part 1 (d)] Include your final MIPS processor schematic in your lab report.

all the blue lines are control input lines



[Part 2 (a.i)] Create a spreadsheet detailing the list of M instructions to be supported in your project alongside their binary opcodes and funct fields, if applicable. Create a separate column for each binary bit. Inside this spreadsheet, create a new column for the N control signals needed by your datapath implementation. The end result should be an $N \times M$ table where each row corresponds to the output of the control logic module for a given instruction.

| 1 | 2 | Instruction | code (Bina | Funct (Binary) | Control Signals | | | | | | | | | | | |
|----|---|-------------|------------|----------------|-----------------|---------|----------|--------|----------|----------|---------|----------|--------|------|--------|---------|
| | | | | | JumpReg | JumpLnk | Unsigned | ALUSrc | MemtoReg | MemWrite | MemRead | RegWrite | RegDst | Jump | Branch | ALUOp |
| 3 | | add | "000000" | "100000" | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | "00000" |
| 4 | | addu | "000000" | "100001" | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | "00001" |
| 5 | | and | "000000" | "100100" | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | "00010" |
| 6 | | not | "000000" | "011111" | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | "00011" |
| 7 | | nor | "000000" | "100111" | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | "00100" |
| 8 | | xor | "000000" | "100110" | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | "00101" |
| 9 | | or | "000000" | "100101" | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | "00110" |
| 10 | | slt | "000000" | "101010" | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | "00111" |
| 11 | | sll | "000000" | "000000" | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | "01000" |
| 12 | | srl | "000000" | "000010" | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | "01001" |
| 13 | | sra | "000000" | "000011" | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | "01010" |
| 14 | | sub | "000000" | "100010" | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | "01011" |
| 15 | | jr | "000000" | "001000" | 1 | 0 | 0 | 1' | 0 | 0 | 0 | 1 | 0 | 1 | 0 | "01100" |
| 16 | | movn | "000000" | "001011" | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | "01101" |
| 17 | | Subu | "000000" | "100011" | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | "11011" |
| 18 | | j | "000010" | "-----" | 0 | 0 | 0 | 1' | 0 | 0 | 0' | 0 | x | 1 | 0 | "01110" |
| 19 | | jal | "000011" | "-----" | 0 | 1 | 0 | 1' | 0 | 0 | 0' | 1 | 1 | 1 | 0 | "01111" |
| 20 | | beq | "000100" | "-----" | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | x | 0 | 1 | "10000" |
| 21 | | addi | "001000" | "-----" | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | "10010" |
| 22 | | repl.qb | "010100" | "-----" | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | "10011" |
| 23 | | addiu | "001001" | "-----" | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | "10100" |
| 24 | | slti | "001010" | "-----" | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | "10101" |
| 25 | | bne | "000101" | "-----" | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | x | 0 | 1 | "10001" |
| 26 | | ori | "001101" | "-----" | 0 | 0 | 1 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | "10110" |
| 27 | | xori | "001110" | "-----" | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | "10111" |
| 28 | | lui | "001111" | "-----" | 0 | 0 | 0 | 1 | 1 | 0 | 1 | 1 | 0 | 0 | 0 | "11000" |
| 29 | | lw | "100011" | "-----" | 0 | 0 | 0 | 1 | 1 | 0 | 1 | 1 | 0 | 0 | 0 | "11001" |
| 30 | | sw | "101011" | "-----" | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 0 | x | 0 | 0 | "11010" |
| 31 | | andi | "001100" | "-----" | 0 | 0 | 1 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | "11100" |

[Part 2 (a.ii)] Implement the control logic module using whatever method and coding style you prefer. Create a testbench to test this module individually, and show that your output matches the expected control signals from problem 1(a).

| | | | | | | | | | |
|-----------------|----------|------|------|--|------|--|------|--|--|
| i_opCode | 6'h00 | 00 | | | | | | | |
| i_funCode | 6'h0C | 00 | 05 | | 0C | | 09 | | |
| o_controlOutput | 13'h110D | 0180 | 0185 | | 110D | | 0189 | | |
| s_OGOutput | 14'h2000 | 2000 | | | | | | | |
| s_RtypeOutput | 13'h110D | 0180 | 0185 | | 110D | | 0189 | | |
| s_RtypeEn | 0 | | | | | | | | |

| | | | | | | | |
|------|--|------|--|------|--|------|--|
| 01 | | 04 | | 0A | | 06 | |
| 00 | | 12 | | 3F | | 00 | |
| 104E | | 0031 | | 1117 | | 1193 | |
| 104E | | 0031 | | 1117 | | 1193 | |
| 0180 | | | | | | | |
| | | | | | | | |

Testing these opcodes comes out with the right values, as shown later, testing with the full processor shows the rest of the values not tested here work correctly

[Part 2 (b.i)] What are the control flow possibilities that your instruction fetch logic must support? Describe these possibilities as a function of the different control flow-related instructions you are required to implement.

BNE: must take a branch not equal input from the control unit and and gate it with NOT ZERO in order to follow with the normal incremented address or change to the branch address

BEQ: must take a branch is equal input from the control unit and and gate with ZERO to choose whether to branch to branch address or stay on normal incremented address


Jal: Jump and link must set register 31 to the PC address incremented by 4. This can be done with a mux on the write address of the register file and porting the pc address +4 into a mux to choose 1 output from mem or pc address when Jal = 1

J: Jump runs to a mux to choose between the jump address and or keep the normal address


Jr: Jump register jumps to the saved register address in RS (31) or the normal address + 4

[Part 2 (b.ii)] Draw a schematic for the instruction fetch logic and any other datapath modifications needed for control flow instructions. What additional control signals are needed?


We needed J, JAL, JR, BNE, and BEQ implementation.




BNE expecting address 7 or 1C



J expecting address output 0000



Jal expecting address 3 with writeDST 31



Jr expecting address 1 again

| | | | |
|----------|----------|--|--|
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |
| 00000004 | | | |
| | | | |
| 00400004 | | | |
| 00000000 | | | |
| 0000000C | | | |
| 01 | | | |
| | | | |
| | | | |
| 0000000C | 00400004 | | |

[Part 2 (c.i.1)] Describe the difference between logical (`srl`) and arithmetic (`sra`) shifts. Why does MIPS not have a `sla` instruction?

A right shift logical simply shifts each bit one position to the right, with the least significant bit being discarded and the most significant bit filled with 0. A left shift logical shifts each bit one to the left, with the least significant bit filled with 0 and the most significant bit being discarded. A right arithmetic shift still shifts each bit one to the right, with the least significant bit being discarded but the most significant bit stays the same, preserving the sign of the number. MIPS does not have an `sla` option because arithmetic and non-arithmetic operations can already be done with `sll`, so it would be the same thing.

[Part 2 (c.i.2)] In your writeup, briefly describe how your VHDL code implements both the arithmetic and logical shifting operations.

For logical shifts, it just pads the shifted bits with 0s, not caring about what the sign bit actually is. For arithmetic operations, it checks whether the sign bit is 1 or 0 (bit 31), and saves it for when it needs to pad the shifted bits at the end.

[Part 2 (c.i.3)] In your writeup, explain how the right barrel shifter above can be enhanced to also support left shifting operations.

If you reverse the input, shift right, and then reverse the output, it's effectively the same thing as shifting left by the same amount. So, instead of coming up with a completely new shifter design, you can just reverse the input and output when a left shift is selected.

[Part 2 (c.i.4)] Describe how the execution of the different shifting operations corresponds to the Modelsim waveforms in your writeup.

Instead of using one barrel shifter file for all of the shifting instructions, we decided to split it up into each shifting operation having a different file. This made more sense to us when implementing the ALU op code and differentiating which shift is done. It also simplifies some of the control signals. The three files are very similar but with slight differences depending on if it was for a left shift or arithmetic shift.

srl:

| | | | | |
|-----------------------------|----------------|-------------------------------------|----------------------------------|---------------------------------|
| /shiftrightlogical/i_In | 00011110010... | (0001001000110100010101011001111000 | 11111111111111111111111111111111 | 0001111001000101100011110010111 |
| /shiftrightlogical/i_shamt | 00000 | 00111 | 11111 | 00000 |
| /shiftrightlogical/o_Out | 00011110010... | (00000000001001000110100010101100 | 00000000000000000000000000000001 | 0001111001000101100011110010111 |
| /shiftrightlogical/s_shift1 | 00011110010... | (00001001000110100010101100111100 | 01111111111111111111111111111111 | 0001111001000101100011110010111 |
| /shiftrightlogical/s_shift2 | 00011110010... | (00000010010001101000101011001111 | 00011111111111111111111111111111 | 0001111001000101100011110010111 |
| /shiftrightlogical/s_shift4 | 00011110010... | (00000000001001000110100010101100 | 00000001111111111111111111111111 | 0001111001000101100011110010111 |
| /shiftrightlogical/s_shift8 | 00011110010... | (00000000001001000110100010101100 | 00000000000000001111111111111111 | 0001111001000101100011110010111 |

This waveform shows the srl instruction. The first test was shifting by 7 bits, which it does correctly. The second test was shifting the max of 31 times, leaving just one bit at the end. The last test was performing no shift. In all cases, the sign bit is never preserved, and the shifted bits are filled with 0, as expected of srl.

sra:

| | | | | |
|--------------------------------|----------------|------------------------------------|----------------------------------|----------------------------------|
| /shiftrightarithmetic/i_In | 00010010001... | (10000001001000110100010101100111 | 11111111111111111111111111111111 | 00010010001101000101011001111000 |
| /shiftrightarithmetic/i_shamt | 00111 | 00111 | 11111 | 00111 |
| /shiftrightarithmetic/o_Out | 00000000001... | (11111111000000100100011010001010 | 11111111111111111111111111111111 | 00000000001001000110100010101100 |
| /shiftrightarithmetic/s_shift1 | 00001001000... | (11000000100100011010001010110011 | 11111111111111111111111111111111 | 00001001000110100010101100111100 |
| /shiftrightarithmetic/s_shift2 | 00000010010... | (11110000001001000110100010101100 | 11111111111111111111111111111111 | 00000010010001101000101011001111 |
| /shiftrightarithmetic/s_shift4 | 00000000001... | (111111111000000100100011010001010 | 11111111111111111111111111111111 | 00000000001001000110100010101100 |
| /shiftrightarithmetic/s_shift8 | 00000000001... | (111111111000000100100011010001010 | 11111111111111111111111111111111 | 00000000001001000110100010101100 |

This waveform shows the sra instruction. The first test was shifting a number with a sign by 7. As you can see, the sign is preserved, effectively dividing the number by 7. The second test shifts 31 times but preserves the sign, so no change is made. Lastly, a number is shifted by 7 that does not have a sign, so no sign bit needs to be preserved.

sll:


| | | | | |
|-------------------------------|----------------|------------------------------------|----------------------------------|----------------------------------|
| /shiftleftlogical/i_In | 00010011010... | (00010010001101000101011001111000 | 11111111111111111111111111111111 | 00010011010001010000111011110100 |
| /shiftleftlogical/i_shamt | 10000 | (00111 | 11111 | 10000 |
| /shiftleftlogical/o_Out | 00001110111... | (00011010001010110011110000000000 | 10000000000000000000000000000000 | 00001110111101000000000000000000 |
| /shiftleftlogical/s_shift1 | 00101111011... | (00001111001101010001011000100100 | 01111111111111111111111111111111 | 00101111011100001010001011001000 |
| /shiftleftlogical/s_shift2 | 00101111011... | (00000011110011010100010110001001 | 00011111111111111111111111111111 | 00101111011100001010001011001000 |
| /shiftleftlogical/s_shift4 | 00101111011... | (000000000001111001101010001011000 | 00000001111111111111111111111111 | 00101111011100001010001011001000 |
| /shiftleftlogical/s_shift8 | 00101111011... | (000000000001111001101010001011000 | 00000000000000011111111111111111 | 00101111011100001010001011001000 |
| /shiftleftlogical/s_reverseIn | 00101111011... | (00011110011010100010110001001000 | 11111111111111111111111111111111 | 00101111011100001010001011001000 |
| /shiftleftlogical/s_Out | 00000000000... | (000000000001111001101010001011000 | 0000000000000000000000000001 | 000000000000000010111101110000 |

This is the waveform for the sll instruction. The first test shifts a random number by 7. The second test shifts the number 32 times, leaving just the most significant bit. The last test shifts a random number 16 times. As you can see, all the shifted numbers are filled with 0 which is expected of sll.

[Part 2 (c.ii.1)] In your writeup, briefly describe your design approach, including any resources you used to choose or implement the design. Include at least one design decision you had to make.

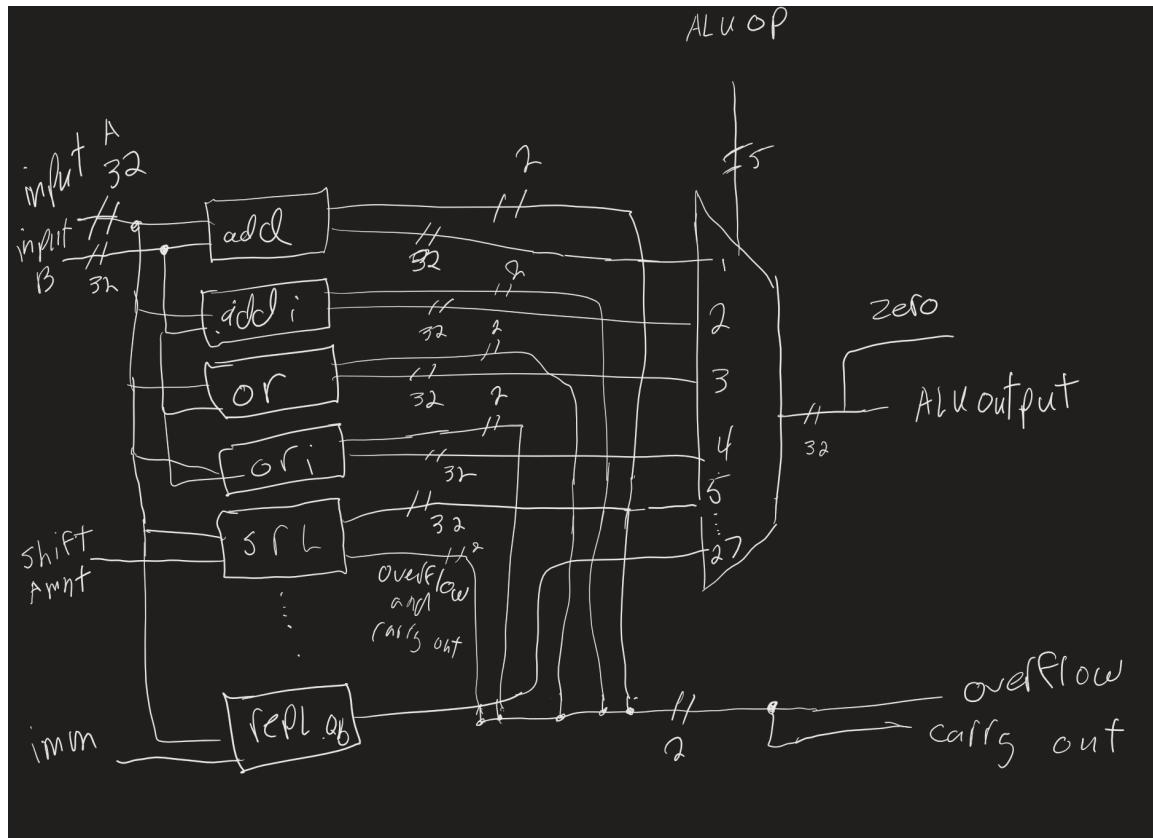
As mentioned before, instead of having our top level ALU schematic contain the shifter, adder, sign extender, and other functional units and then creating signals that choose what operation was done, we decided to make a schematic for every operation, even if they used the same functional unit. For example, srl, sll, and sra all have their own files. Obviously, they are very similar to each other but with a few differences. This was somewhat redundant since many operations are similar, but it didn't take much longer and made it a lot easier for us to understand which operation is chosen via the ALU opcode. Because of this, we also had to make a new input specifically for the repl instruction that is only active when the repl ALU op is chosen. We also had to create a 29-1 Multiplexor that the ALU Op Code chooses which instruction to take. Depending on the op code, that instruction is put through the multiplexer and only its output is outputted from the ALU

[Part 2 (c.ii.2)] Describe how the execution of the different operations corresponds to the Modelsim waveforms in your writeup.




This waveform was a test for the 29 to 1 - 32 bit Multiplexor. This multiplexor is used by the ALU and takes in the ALUOpCode control bus from the control unit. This control bus will choose the output of the ALU as well as choose the overflow and carry-out bits. The above waveform is acting correctly. It is a bit difficult to explain what's going on in the waveform, but the testbench is in the ALU folder if you wanna look at it. All instructions will run as if they were called, but only the correct instructions output will come out of the ALU due to this MUX and the ALUOpCode control signal.

[Part 2 (c.iii)] Draw a simplified, high-level schematic for the 32-bit ALU. Consider the following questions: how is Overflow calculated? How is Zero calculated? How is `slt` implemented?




Overflow is calculated by taking the carry out bit and second to last carry out bit and XORing them. The zero is calculated by making the least significant bit of the ALU output either 1 or 0. We do this instead of making a whole different output in order to make the ALU outputs simple, all we need to do is correctly setup the bit to the right place in the top level processor. SLT is implemented by subtracting the `rt` register from `rs` and if the output is negative, then `rs` is smaller than `rt` and therefore the output of this instruction is `0x00000001`. In the opposite case where the subtraction output is positive, then the output of the instruction is `0x00000000`. That last bit is used by the branch instructions.


[Part 2 (c.v)] Describe how the execution of the different operations corresponds to the Modelsim waveforms in your writeup.



The first two tests are for **add (ALU op 0)** instruction. They both work as intended and calculate overflow as intended. The second two tests are for **addu (ALU op 1)**, they work the exact same as add but they don't care about overflow. The next two tests are for **sub (ALU op 11)** instruction. They both subtract the two inputs and make sure to keep the overflow. ShiftImm and replImm signals are always don't care unless we are shifting or replicating.



The next two tests are for **subu (ALU op -5)**, it works the same as sub but just doesn't care about overflow. The next test was for **not (ALU op 3)**, it simply inverts all of the bits, which in decimal would flip the sign and then subtract one as the output shows. The last three tests are for testing **slt (ALU op 7)**, which produces a 1 when RS is less than RT; this also includes the case in which they are both 0, then the output is still 0.



The next tests were for **and (ALU op 2)** instruction and **or (ALU op 6)** instruction, these cases are trivial but work perfectly.



The next tests here are for **xor (ALU op 5)** and **nor (ALU op 4)**, these are also trivial and work as expected even with edge cases such as all 0s or all 1s.

[illegible]

Note that in this step, we did not test the immediate instructions because the bit extender is not technically part of the ALU, and they function the exact same as their non-immediate counterparts once extended.

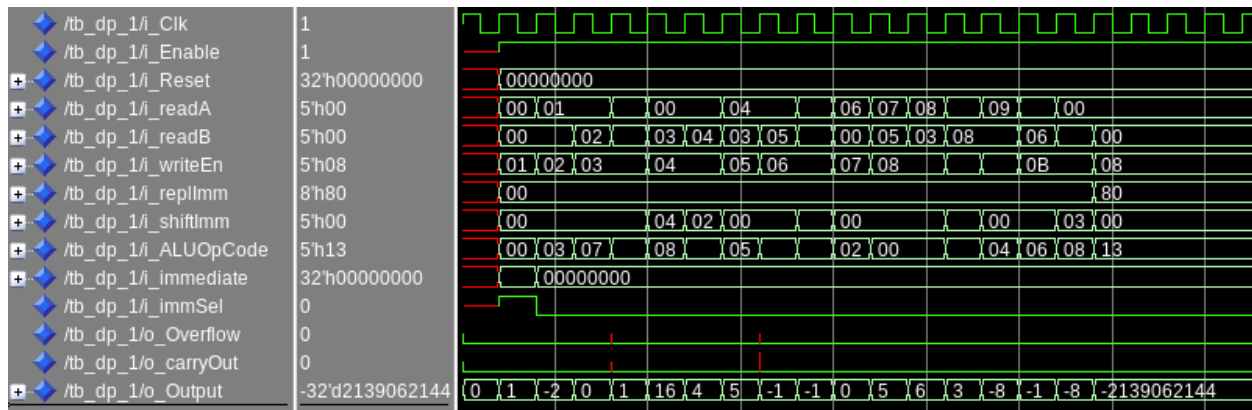
```

---ALU dataflow using Lab 2 datapath---

addi $1, $0, 1           # Place "1" in $1
not $2, $1                # negate "1" and place into 2 ($2 is now -2)
slt $3, $1, $2            # check if $1 is less than $2 and place into $3 ($3 is now 0)
add $3, $3, $1            # Place "1" in $3 ($3 is now 1)
sll $4, $3, 4             # shift $3 left 4 times, place into $4 ($4 should now be 16)
sra $4, $4, 2             # Shift $4 right 2 and place into 4 ($4 should now be 4)
xor $5, $4, $3            # xor $3 and $4 and place into $5 ($5 should have the value 5)
sub $6, $4, $5            # subtract $4 from $5 ($6 should be -1)
sra $6, $6, 12            # shift right arithmetic $6 12 times ($6 should still be -1)
and $7, $6, 0            # and $6 with 0 and place into $7 ($7 should be 0)
add $8, $7, $5            # add the value of $5 and $7 into $8 ($8 should have value of 5)
add $8, $8, $3            # add $3 to $8 and store in $8 ($8 should have 6)
srl $9, $8, 1            # divide $8 by 2 and store in $9 ($9 should contain 3)
nor $10, $9, $8          # nor $8 and $9 ($10 should contain -8)
or $11, $10, $6          # or $10 with $6 and store in $11 ($11 should be -1)
sll $11, $11, 3          # shift left 3 times ($8 should finished with a value of -8)
repl $12, 01000000      #replicate 010000000 and place into $12.

final value of $12 should be 0x80808080 (-2139062144 in decimal) and the value of 11 should be -8 after everything

```



Here is the waveform for our test. The signal o_Output contains the output after each instruction so you can follow it along to see how it changes. This also contains all of the other intermediate signals used to perform the operations.

```

1  addi $t1, $0, 1      # Place "1" in $1
2  not  $t2, $t1        # negate "1" and place into 2 ($2 is now -2)
3  slt  $t3, $t1, $t2   # check if $1 is less than $2 and place into $3 ($3 is now 0)
4  add  $t3, $t3, $t1   # Place "$1" in $3 ($3 is now 1)
5  sll  $t4, $t3, 4     # shift $3 left 4 times, place into $4 ($4 should now be 16)
6  sra  $t4, $t4, 2     # Shift $4 right 2 and place into 4 ($4 should now be 4)
7  xor  $t5, $t4, $t3   # xor $3 and $4 and place into $5 ($5 should have the value 5)
8  sub  $t6, $t4, $t5   # subtract $4 from $5 ($6 should be -1)
9  sra  $t6, $t6, 12    # shift right arithmetic $6 12 times ($6 should still be -1)
10 and  $t7, $t6, 0     # and $6 with 0 and place into $7 ($7 should be 0)
11 add  $s0, $7, $t5    # add the value of $5 and $7 into $8 ($8 should have value of 5)
12 add  $s0, $s0, $t3   # add $3 to $8 and store in $8 ($8 should have 6)
13 srl  $s1, $s0, 1     # divide $8 by 2 and store in $9 ($9 should contain 3)
14 nor  $s2, $s1, $s0   # nor $8 and $9 ($10 should contain -8)
15 or   $s3, $s2, $t6   # or $10 with $6 and store in $11 ($11 should be -1)
16 sll  $s3, $s3, 3     # shift left 3 times ($8 should finished with a value of -8)
17

```

| | | |
|--------|----|------------|
| \$zero | 0 | 0x00000000 |
| \$at | 1 | 0x00000000 |
| \$v0 | 2 | 0x00000000 |
| \$v1 | 3 | 0x00000000 |
| \$a0 | 4 | 0x00000000 |
| \$a1 | 5 | 0x00000000 |
| \$a2 | 6 | 0x00000000 |
| \$a3 | 7 | 0x00000000 |
| \$t0 | 8 | 0x00000000 |
| \$t1 | 9 | 0x00000001 |
| \$t2 | 10 | 0xffffffff |
| \$t3 | 11 | 0x00000001 |
| \$t4 | 12 | 0x00000004 |
| \$t5 | 13 | 0x00000005 |
| \$t6 | 14 | 0xffffffff |
| \$t7 | 15 | 0x00000000 |
| \$s0 | 16 | 0x00000006 |
| \$s1 | 17 | 0x00000003 |
| \$s2 | 18 | 0xffffffff |
| \$s3 | 19 | 0xffffffff |

Here is our program ran in MARS. \$S3 contains the output of \$11 and finishes with a value -8 as expected. All of the other registers contain the correct values throughout the program (did not test repl on MARS).

[Part 3] In your writeup, show the Modelsim output for each of the following tests, and provide a discussion of result correctness. It may be helpful to also annotate the waveforms directly.

[Part 3 (a)] Create a test application that makes use of every required arithmetic/logical instruction at least once. The application need not perform any particularly useful task, but it should demonstrate the full functionality of the processor (e.g., sequences of many instructions executed sequentially, 1 per cycle while data written into registers can be effectively retrieved and used by later instructions). Name this file Proj1_base_test.s.

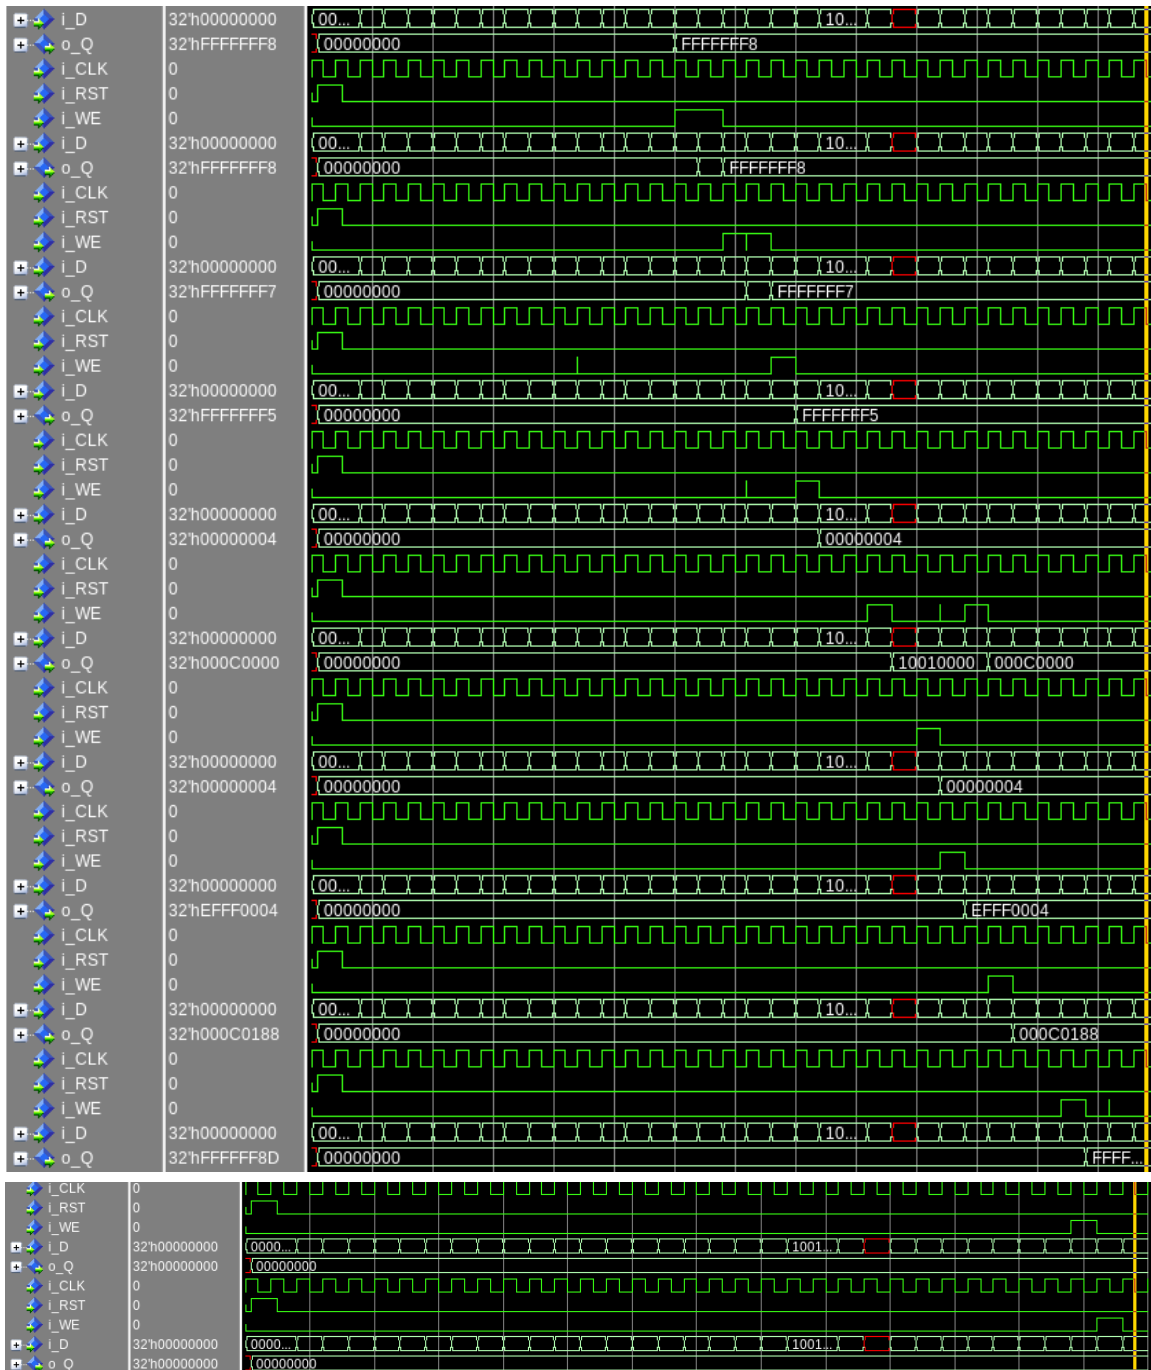
```
addi $1, $0, 1      # Place "1" in $1
not $2, $1          # negate "1" and place into 2 ($2 is now -2)
slt $3, $1, $2      # check if $1 is less than $2 and place into $3 ($3 is now 0)
add $3, $3, $1       # Place "1" in $3 ($3 is now 1)
sll $4, $3, 4        # shift $3 left 4 times, place into $4 ($4 should now be 16)
sra $4, $4, 2        # Shift $4 right 2 and place into 4 ($4 should now be 4)
xor $5, $4, $3       # xor $3 and $4 and place into $5 ($5 should have the value 5)
sub $6, $4, $5       # subtract $4 from $5 ($6 should be -1)
sra $6, $6, 12       # shift right arithmetic $6 12 times ($6 should still be -1)
and $7, $6, 0        # and $6 with 0 and place into $7 ($7 should be 0)
add $8, $7, $5       # add the value of $5 and $7 into $8 ($8 should have value of 5)
add $8, $8, $3       # add $3 to $8 and store in $8 ($8 should have 6)
srl $9, $8, 1        # divide $8 by 2 and store in $9 ($9 should contain 3)
nor $10, $9, $8      # nor $8 and $9 ($10 should contain -8)
or $11, $10, $6      # or $10 with $6 and store in $11 ($11 should be -1)
sll $11, $11, 3      # shift left 3 times ($11 should finished with a value of -8)
repl.qb $12, 29      #replicate 010000000 and place into $12.
addu $12, $11, $6
addiu $13, $12, -2
andi $14, $13, 6
addi $15, $15, 0x10010000
sw $14, 0($15)
lw $16, 0($15)
subu $17, $16, $15
lui $15, 12
xori $18, $15, 392
ori $19, $18, -123
slti $20, $18, 0
movn $21, $20, $19
halt
```

This is the Proj1_base_test.s file in Mars. It uses every required arithmetic/logical instruction at least once. It does not do anything important, but demonstrates the ability of the processor to perform instructions sequentially.

| Name | Number | Value |
|--------|--------|--------------|
| \$zero | 0 | 0x00000000 |
| \$at | 1 | 0xffffffff85 |
| \$v0 | 2 | 0xfffffffffe |
| \$v1 | 3 | 0x00000001 |
| \$a0 | 4 | 0x00000004 |
| \$a1 | 5 | 0x00000005 |
| \$a2 | 6 | 0xffffffffff |
| \$a3 | 7 | 0x00000000 |
| \$t0 | 8 | 0x00000006 |
| \$t1 | 9 | 0x00000003 |
| \$t2 | 10 | 0xfffffffff8 |
| \$t3 | 11 | 0xfffffffff8 |
| \$t4 | 12 | 0xfffffffff7 |
| \$t5 | 13 | 0xfffffffff5 |
| \$t6 | 14 | 0x00000004 |
| \$t7 | 15 | 0x000c0000 |
| \$s0 | 16 | 0x00000004 |
| \$s1 | 17 | 0xefff0004 |
| \$s2 | 18 | 0x000c0188 |
| \$s3 | 19 | 0xffffffff8d |
| \$s4 | 20 | 0x00000000 |
| \$s5 | 21 | 0x00000000 |
| \$s6 | 22 | 0x00000000 |
| \$s7 | 23 | 0x00000000 |
| \$t8 | 24 | 0x00000000 |
| \$t9 | 25 | 0x00000000 |
| \$k0 | 26 | 0x00000000 |
| \$k1 | 27 | 0x00000000 |
| \$gp | 28 | 0x10008000 |
| \$sp | 29 | 0x7ffffc |
| \$fp | 30 | 0x00000000 |
| \$ra | 31 | 0x00000000 |
| pc | | 0x00400088 |
| hi | | 0x00000000 |
| lo | | 0x00000000 |

These are the register outputs in Mars after running the previous program. Every register has its intended value.





Here are the output waveforms generated by the tool flow tests. As you can see, each register contains the same values as it does in the Mars simulation ranging from register \$1 to \$21, indicating the processor works correctly.

[Part 3 (b)] Create and test an application which uses each of the required control-flow instructions and .data (i.e., the number of activation records on the stack is at least 4). Name this file Proj1_cf_test.s.

```
main:
    lui $sp, 0x7FFF
    ori $sp, 0xEFF4
    addi $a0, $0, 5 # set n = 5
    jal recursive
    addi $t9, $0, 55
    bne $v0, $t9, failure
    j exit

recursive:
    addi $sp, $sp, -8      # space for two words
    sw $ra, 4($sp)        # save return address
    sw $a0, 0($sp)        # temporary variable to hold n
    li $v0, 1
    slti $t0, $a0, 1
    bne $t0, $0, recexit
    addi $a0, $a0, -1
    jal recursive
    lw $a0, 0($sp)        # retrieve original n
    add $v0, $v0, $a0     # n + ((n - 1) + (n - 2) + ... (n - n))

recexit:
    lw $ra, 4($sp)        # restore $ra
    addi $sp, $sp, 8      # restore $sp
    jr $ra                # back to caller

failure:
    addi $t8, $0, 1
    j exit

exit:
    li $v0, 5
    halt
```

This is a script that takes a value n and calculates n! recursively.

| | | | | | | | | | | | |
|----------|----------|----------|----------|----------|----------|----------|----------|--|----------|--|----------|
| 7FFF0000 | | | 7FFFEFF4 | | 00000005 | | 00400010 | | 7FFFEFEC | | XXXXXXXX |
| XXXXXXXX | 00000001 | 00000000 | 00000000 | 00000004 | 0040003C | 7FFFEFE4 | XXXXXXXX | | | | |
| XXXX... | 00000001 | 00000000 | 00000000 | 00000003 | 0040003C | 7FFFEFDC | XXXXXXXX | | | | |
| | 00000001 | 00000000 | 00000000 | 00000002 | 0040003C | 7FFFEFD4 | XXXX | | | | |

| | | | | | | | |
|----|----------|----------|----------|----------|----------|----------|-----|
| | 00000001 | 00000000 | 00000000 | 00000001 | 0040003C | 7FFFEFC4 | XXX |
| XX | 00000001 | 00000000 | 00000000 | 00000000 | 0040003C | 7FFFEFC4 | XX |

[Part 3 (c)] Create and test an application that sorts an array with N elements using the BubbleSort algorithm ([link](#)). Name this file Proj1_bubblesort.s.

```

e sort:

    add $t0,$zero,$zero
    addi $t1, $t1, -1
loop:

    beq $t0,$t1,done

    add $t4,$zero,$zero

    innerLoop:

        beq $t4, $t1, continue
        sll $t5, $t4, 2
        add $t6, $s0, $t5

        lw $s1, 0($t6)
        lw $s2, 4($t6)

        slt $t7, $s2, $s1

        beq $t7, $0, good

        sw $s2, 0($t6)
        sw $s1, 4($t6)

    good:
        addi $t4, $t4, 1
        j innerLoop

    continue:
        addi $t0, $t0, 1
        j loop

```

This is the sorting algorithm for bubble sort, accessing the numbers in the array from memory and switching them if the first number is larger then the second number.

```

newLine: .ascii "\n" # Newline char
.align 2
numbers: .word 7 3 1 46 24 10 5 34 100 99

    text

```

This is the starting array of numbers for bubble sort. These numbers can be changed to anything and make any size of array and the bubble sort code we wrote will still work.

before:
The array starts at 7

after:
the array starts at 1

[Part 4] report the maximum frequency your processor can run at and determine what your critical path is. Draw this critical path on top of your top-level schematics. What components would you focus on to improve the frequency?

We would probably need to optimize the ALU because no matter what instruction is called, all instructions run in the ALU and only the correct instruction is outputted out of the ALU. So changing that fact would greatly increase our frequency

The red line below is the critical path, the critical path goes through the reg file, the ALU, both memory files, and a couple mux's

